





Parallelizing BLAKE3




Rehan Vipin & Samarth Mathur
Team - 6



Recap

BLAKE3 is the fastest cryptographic hash function (throughput - 5 GB/s). It has security equivalent to (or better than) existing standards and is being widely adopted.

It has a merkle tree structure, allowing unbounded parallelism. ->



At the end of phase-1 we had a serial-algorithm based version ready. It was a simple but a complete port of the BLAKE3 implementation.

We discussed multiple approaches available.

In Phase-2 we :

- Implemented $\frac{3}{4}$ of the them
- Design+Implemented a new algorithm
- Benchmarked, documented and examined, in-detail, every version we made.

10 MB/s -> 1.2 GB/s

4 Different techniques

- OpenMP
- SIMD (AVX2)
- GPU (Dark)
- GPU (Dynamic parallelism)

Before all of that, we needed a base -

A new technique in BLAKE3 hashing, parallelize wherever possible. We call this BLAZE3. Explained through diagrams from the notebook:

Rest through code and demo

Phase-2 Timeline

1. Design and implement new parallel algorithm
2. Implement openmp version and testing framework
3. Optimize OMP, profile, collect many tests
4. GPU - set up base (a & b done in parallel)
 - a. Dark version
 - b. Dynamic parallelism version
5. Profile, identify hotspots in GPU version, improve it
6. Implement SIMD v1, v2 ...

Optimizations, major and minor, were being done continuously.

OpenMP

Needed a fork-join model with nesting.

Used tasks in openmp with 2 threads and 3 levels of nesting.

Heavy on CPU - ~80%. Tried to set up enhancers -

<https://github.com/pmodels/bolt> but documentation is horrible. Async too slow.

Designed to be closely related to serial version - just use the ``-fopenmp`` flag.

More in benchmarks...

GPU (Dynamic Parallelism)

Will show shortly.

GPU (Dark)

Same `hash_many` interface, but internally uses a new `dark hash` technique.

Similar to parallel-reduction, but not exactly!

Tried many techniques for improvement:

- Shared memory, reduce local memory ...
- Algorithmic improvements to minimize kernels & divergent code
- Thrust and pinned memory
- Zero-Copy Registered memory -> slowest!

More soon...

SIMD

AVX2 provides access to 128 & 256 bit registers to speed up core functions.
AVX-512 is not yet widely used.

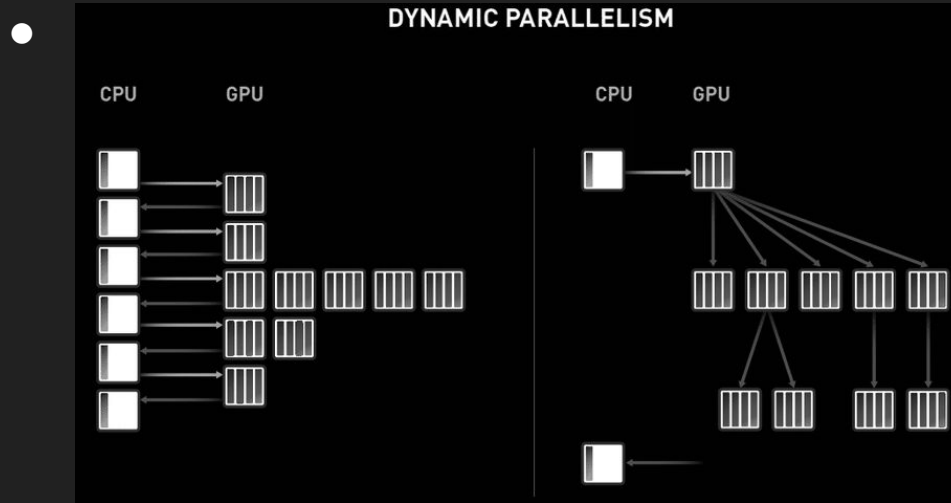
Activated only along with OpenMP. Together, they give us the fastest throughput.

Can also be used in single-thread mode but technically a “parallelism”, therefore separated.

Not available in GPUs :(


Dynamic Parallelism

- It is an extension to the CUDA programming model enabling a CUDA kernel to create new thread grids by launching new kernels.



- Used in algorithms like “The Mandelbrot Set”, essentially algorithms where recursive nature can’t be avoided.

How did we try it?




- A divide and conquer approach, somewhat like merge sort.
- Kernels launching recursively for every half.
- We split from the chunk in the middle
- Code walkthrough
- Unfortunately, it was slower and also incorrect.

Where did we fail?

- Intense testing and debugging for the past one week.
- So what all did we try to mitigate this?
 - `d_actual_compress` , a kernel in our implementation was passing individual test cases but was failing inside the program itself
 - Wrote individual test cases of each of these function
 - Wrote individual tests for the smaller kernels
 - In conclusion we found that `d_actual_compress` was indeed receiving the right values but was “returning” wrong ones, even though its individual case was passing and the values computed inside were right.
 - A detailed `.cu` file of this test can be found in `test_in` directory
- But, the hashes were consistently wrong, so we did do a speed run.

Benchmarks

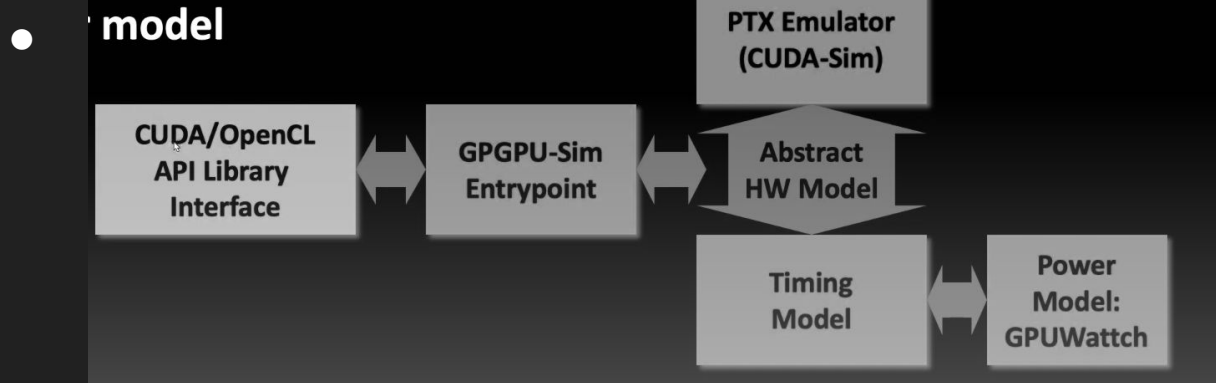


GPGPU-SIM

Optimization at the architectural level, not quite :(

What is GP-GPUSIM

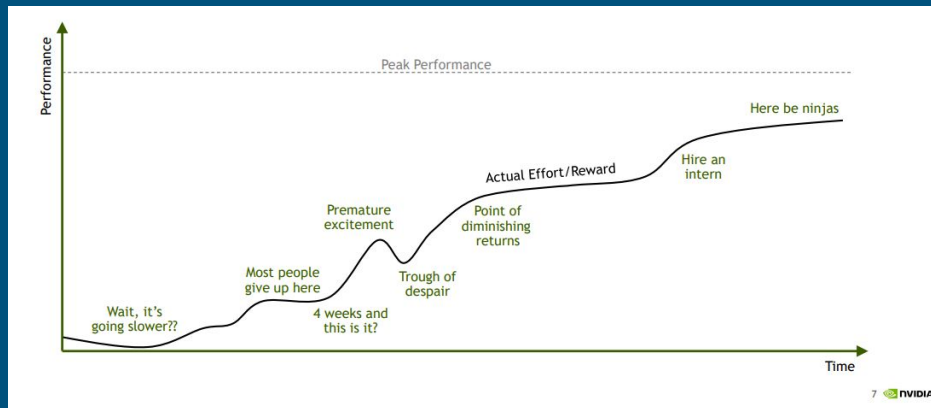
- It is a cycle level simulator that can allow you to “make” architectural changes. You can tweak things at the architectural level by editing a configuration file.
- You can change things at the high level and low level.



Why did this not work?

- Dark Hash uses a PTX instruction called “shf.l.wrap.b32” which is not yet implemented by gp-gpusim.
- The Dynamic parallelism approach did not work because dynamic parallelism throws error for newer versions of CUDA in gp-gpusim.
- The PTX instruction could’ve been implemented but we were short on time as it required a good amount of documentation and codebase exploration. Programming gp-gpusim for handling Dynamic Parallelism seemed like a task too complex.

Inferences, benchmarks & more...



Primitive implementations:


```
❖ apollo_nox cuda on cpu-thread > python bench.py 512000000 og
File size in bytes: 512000000
8e9a08f9552b774d4c08d5735715da83755a4feec0e6b3c88276cffdbb37e1ef
Execution time: 20.88s
```

Serial

```
❖ apollo_nox cuda on cpu-thread > python bench.py 512000000 og
File size in bytes: 512000000
8e9a08f9552b774d4c08d5735715da83755a4feec0e6b3c88276cffdbb37e1ef
Execution time: 9.13s
```

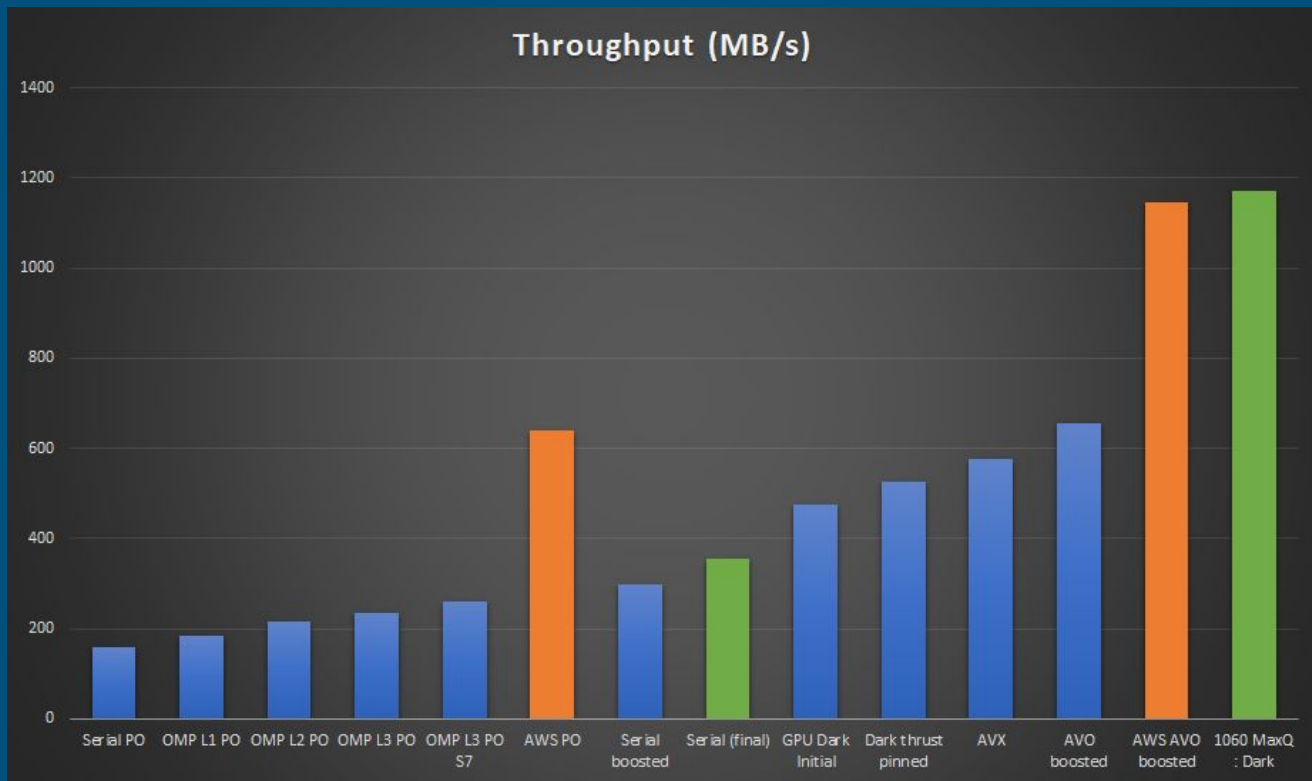
Two level nesting OMP

```
❖ apollo_nox cuda on cpu-thread > clang++ main.cpp -fopenmp
❖ apollo_nox cuda on cpu-thread > python test.py
Running test cases
35 test cases passed out of 35
❖ apollo_nox cuda on cpu-thread > python bench.py 512000000 og
File size in bytes: 512000000
8e9a08f9552b774d4c08d5735715da83755a4feec0e6b3c88276cffdbb37e1ef
Execution time: 7.70s
8e9a08f9552b774d4c08d5735715da83755a4feec0e6b3c88276cffdbb37e1ef
Original blake3 execution time: 0.13s
❖ apollo_nox cuda on cpu-thread > python bench.py 512000000 og
File size in bytes: 512000000
```



3 level nesting OMP

Long story short



Execution on AWS t2.xlarge

```
ubuntu@ip-172-31-87-198:~/zipp/openmp$  
File size in bytes: 4000000000  
17de18ad8a9bfc16b082258537d843681d754b2  
Execution time: 3.53s  
17de18ad8a9bfc16b082258537d843681d754b2  
Sequential execution time: 16.02s  
17de18ad8a9bfc16b082258537d843681d754b2  
Original blake3 execution time: 0.30s  
ubuntu@ip-172-31-87-198:~/zipp/openmp$  
File size in bytes: 4000000000  
17de18ad8a9bfc16b082258537d843681d754b2  
Execution time: 3.49s  
17de18ad8a9bfc16b082258537d843681d754b2  
Sequential execution time: 15.92s  
17de18ad8a9bfc16b082258537d843681d754b2  
Original blake3 execution time: 0.30s  
ubuntu@ip-172-31-87-198:~/zipp/openmp$  
File size in bytes: 4000000000  
17de18ad8a9bfc16b082258537d843681d754b2  
Execution time: 3.45s  
17de18ad8a9bfc16b082258537d843681d754b2  
Sequential execution time: 15.61s  
17de18ad8a9bfc16b082258537d843681d754b2  
Original blake3 execution time: 0.30s
```

```
1 [ 0.0%] 5 [ 0.0%]  
2 [ 0.0%] 6 [ 0.0%]  
3 [ 0.7%] 7 [ 0.0%]  
4 [|||||] 8 [ 0.0%]  
Mem[|||||] 235M/31.4G Tasks: 35, 72 thr; 2 running  
Swp[ 0K/0K] Load average: 1.55 1.85 1.09  
Uptime: 00:14:02
```

PID	USER	PRI	NI	VIRT	RES	SHR	S	CPU%	MEM%	TIME+	Command
2962	ubuntu	20	0	8528	5684	2928	R	100.	0.0	0:20.04	bench/a.out bench/bencher.bin
2160	ubuntu	20	0	8248	3996	3144	R	0.7	0.0	0:02.26	htop

```
1 [|||||] 5 [|||||] 100.0%  
2 [|||||] 6 [|||||] 100.0%  
3 [|||||] 7 [|||||] 100.0%  
4 [|||||] 8 [|||||] 100.0%  
Mem[|||||] 238M/31.4G Tasks: 35, 75 thr; 8 running  
Swp[ 0K/0K] Load average: 3.01 0.87 0.29  
Uptime: 00:05:37
```

PID	USER	PRI	NI	VIRT	RES	SHR	S	CPU%	MEM%	TIME+	Command
2321	ubuntu	20	0	512M	4888	4340	R	800.	0.0	0:24.63	./a.out bench/bencher.bin
2323	ubuntu	20	0	512M	4888	4340	R	101.	0.0	0:03.08	./a.out bench/bencher.bin
2324	ubuntu	20	0	512M	4888	4340	R	100.	0.0	0:03.08	./a.out bench/bencher.bin
2322	ubuntu	20	0	512M	4888	4340	R	100.	0.0	0:03.07	./a.out bench/bencher.bin
2327	ubuntu	20	0	512M	4888	4340	R	100.	0.0	0:03.07	./a.out bench/bencher.bin
2325	ubuntu	20	0	512M	4888	4340	R	99.8	0.0	0:03.07	./a.out bench/bencher.bin
2328	ubuntu	20	0	512M	4888	4340	R	99.8	0.0	0:03.07	./a.out bench/bencher.bin
2326	ubuntu	20	0	512M	4888	4340	R	99.1	0.0	0:03.05	./a.out bench/bencher.bin
2160	ubuntu	20	0	8248	3876	3132	R	1.3	0.0	0:00.52	htop

Perf analysis for CPU version


Overhead	Command	Shared Object	Symbol
71.83%	a.out	a.out	[.] compress
10.99%	a.out	a.out	[.] Chunk::compress_chunk
5.48%	a.out	[kernel.kallsyms]	[k] copy_user_enhanced_fast_string
1.47%	a.out	libc-2.31.so	[.] read
1.10%	a.out	a.out	[.] hash_many
1.04%	a.out	libstdc++.so.6.0.28	[.] std::istream::read
0.99%	a.out	libc-2.31.so	[.] 0x000000000018e8fa

Let's move onto the GPU!

CUDA Insights

Dark hash (l)Nsight - v1

43	Function name	Timestamps	Duration	PID
44	cudaMalloc		162.210 μ s	
45	cudaMemcpy		70.093 μ s	
46	h_compute		14.261 μ s	
47	h_compute		3.667 μ s	
48	h_compute		2.430 μ s	
49	h_compute		2.085 μ s	
50	h_compute		2.002 μ s	
51	h_compute		2.073 μ s	
52	h_compute		2.005 μ s	
53	h_compute		2.019 μ s	
54	h_compute		1.963 μ s	
55	h_compute		1.961 μ s	
56	cudaMemcpy		967.043 μ s	
57	cudaFree		200.419 μ s	
58	cudaMalloc		185.205 μ s	




CUDA Insights

Dark Hash (I)Nsight-v2

33	cudaFree	4.53855s	237.203 μ s
34	cudaMalloc	4.54008s	190.663 μ s
35	cudaMemcpy	4.54027s	99.956 μ s
36	cudaDeviceSynchronize	4.54037s	631.416 μ s
37	h_compute	4.5411s	19.394 μ s
38	h_compute	4.54102s	3.756 μ s
39	h_compute	4.54103s	2.457 μ s
40	h_compute	4.54103s	2.242 μ s
41	h_compute	4.54103s	2.186 μ s
42	h_compute	4.54104s	2.167 μ s
43	h_compute	4.54104s	2.101 μ s
44	h_compute	4.54104s	4.719 μ s
45	h_compute	4.54105s	2.345 μ s
46	h_compute	4.54105s	2.007 μ s
47	cudaDeviceSynchronize	4.54105s	386.255 μ s
48	cudaMemcpy	4.54144s	69.536 μ s
49	cudaFree	4.54151s	203.042 μ s
50	cudaMalloc	4.54348s	176.297 μ s

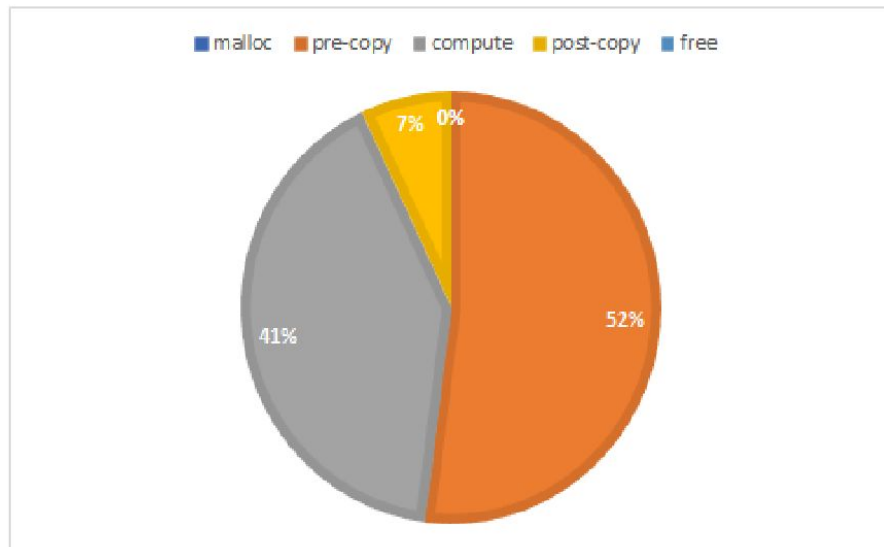
■ malloc ■ pre-copy ■ compute ■ post-copy ■ free



CUDA Insights

Dark hash (I)Nsight v3

86	cudaMemcpy	0.945279s	66.553 μ s
87	cudaMemcpy	0.946224s	59.554 μ s
88	cudaDeviceSynchronize	0.946284s	464.553 μ s
89	h_compute	0.946749s	11.981 μ s
90	h_compute	0.946762s	5.307 μ s
91	h_compute	0.946767s	2.751 μ s
92	h_compute	0.94677s	2.240 μ s
93	h_compute	0.946773s	2.117 μ s
94	h_compute	0.946775s	2.220 μ s
95	h_compute	0.946778s	2.183 μ s
96	h_compute	0.94678s	2.089 μ s
97	h_compute	0.946782s	2.205 μ s
98	h_compute	0.946785s	2.067 μ s
99	cudaDeviceSynchronize	0.946787s	379.058 μ s
100	cudaMemcpy	0.947166s	69.360 μ s
101	cudaMemcpy	0.9481s	56.911 μ s



CUDA Insights

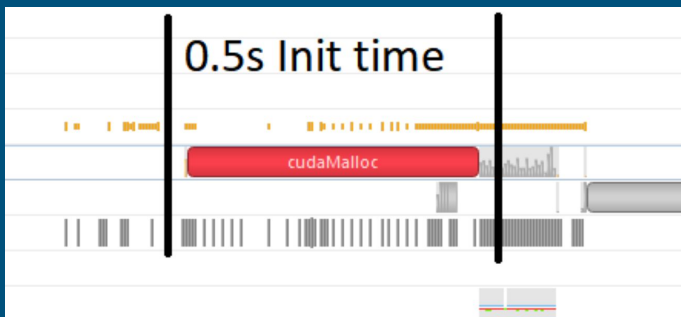
```
==21456== NVPROF is profiling process 21456, command: a.exe bench/bencher.bin
```

```
4610597608768d9d795efdb7cafc5c68ab62616e4d606884cf693ff7ba56d35e
```

```
==21456== Profiling application: a.exe bench/bencher.bin
```

```
==21456== Profiling result:
```

	Type	Time(%)	Time	Calls	Avg	Min	Max	Name
GPU activities:		56.67%	746.98ms	1919	389.26us	1.9840us	408.26us	[CUDA memcpy HtoD]
		43.25%	570.12ms	19136	29.793us	4.7360us	218.30us	h_compute(Chunk*, int, int)
		0.08%	1.0491ms	1919	546ns	320ns	704ns	[CUDA memcpy DtoH]
API calls:		80.75%	1.48715s	3838	387.48us	27.300us	773.40us	cudaMemcpy
		13.82%	254.54ms	1	254.54ms	254.54ms	254.54ms	cudaMalloc
		3.36%	61.887ms	19136	3.2340us	1.9000us	58.600us	cudaLaunchKernel



This one is using thrust + pinned memory, but missing a few optimizations we added later on. Ran on my GPU. My GPU(Low energy) only has a transfer rate of 3Gbps Samarth's GPU(Gaming laptop) can transfer at 12Gbps

What if we use a better GPU?

```
==22247== NVPROF is profiling process 22247, command: ./a.out bench/bencher.bin
9c8ff3575045aa272e21e4d3a512f03908abc6dbd0115dc508210f65adf2022b
==22247== Profiling application: ./a.out bench/bencher.bin
==22247== Profiling result:
```

Type	Time(%)	Time	Calls	Avg	Min	Max	Name
GPU activities:	62.19%	626.73ms	39124	16.018us	6.1120us	637.48us	h compute(Chunk*, int, int)
	37.43%	377.21ms	3915	96.350us	832ns	227.08us	[CUDA memcpy HtoD]
	0.38%	3.8643ms	3915	987ns	800ns	18.849us	[CUDA memcpy DtoH]
API calls:	78.69%	1.00646s	7830	128.54us	9.9240us	779.35us	cudaMemcpy
	12.01%	153.59ms	1	153.59ms	153.59ms	153.59ms	cudaMalloc
	9.04%	115.61ms	39124	2.9550us	2.0810us	499.55us	cudaLaunchKernel
	0.09%	1.1296ms	7	161.37us	4.9390us	607.61us	cudaFreeHost
	0.09%	1.0922ms	7	156.03us	6.8980us	534.36us	cudaMallocHost
	0.04%	450.64us	1	450.64us	450.64us	450.64us	cuDeviceTotalMem
	0.03%	405.22us	101	4.0120us	395ns	180.74us	cuDeviceGetAttribute
	0.01%	143.88us	1	143.88us	143.88us	143.88us	cudaFree
	0.01%	79.739us	1	79.739us	79.739us	79.739us	cuDeviceGetName
	0.00%	34.854us	1	34.854us	34.854us	34.854us	cuDeviceGetPCIBusId
	0.00%	3.9780us	3	1.3260us	558ns	2.6800us	cuDeviceGetCount
	0.00%	2.6630us	2	1.3310us	417ns	2.2460us	cuDeviceGet
	0.00%	2.4340us	7	347ns	175ns	522ns	cudaGetLastError
	0.00%	959ns	1	959ns	959ns	959ns	cuDeviceGetUuid

Overall results

Speed up over sequential version :

- OpenMP + AVX - 4.5
- GPU Dark - 3.5
- GPU Dynamic Parallelism -

End result: Easily 2.5x faster than the inbuilt SHA-256 hash utility on Windows.

Original BLAKE3 is 8x faster than us, but ...

Work split up

Rehan

Contribution -

- Parallel algorithm
- OpenMP
- AVX
- GPU (Dark)

Hours worked -

75

Samarth

Contribution -

- Serial code
- Parallel algorithm
- GPU (Dynamic parallelism)
- GPU Optimizations

Hours worked -

80

Thank you!

